

# Mustafa Ozan Karsavuran

Computer Systems Engineer 2, Lawrence Berkeley National Laboratory, Berkeley, California, 94720 USA  
✉ ozan.karsavuran@gmail.com | ✉ MOKarsavuran@lbl.gov | ✉ ozan-karsavuran-421a7419

## SUMMARY

---

Experienced researcher specializing in high-performance computing, parallel algorithms, and sparse matrix/tensor computations. Skilled in MPI and OpenMP programming, GPU acceleration, and hypergraph partitioning models. Proven experience in performance optimization, algorithm development, and successful collaboration on large-scale scientific computing projects.

## EDUCATION

---

### Bilkent University, Department of Computer Engineering, Turkey

- **Doctor of Philosophy** GPA: 3.47/4.00 September 2014 – August 2020
  - Thesis: Reducing Communication Overhead in Sparse Matrix and Tensor Computations
  - Advisor: Prof. Dr. Cevdet Aykanat
- **Master of Science** GPA: 3.60/4.00 September 2012 – September 2014
  - Thesis: Increasing Data Reuse in Parallel SpMV and SpMTM Multiply on Shared-Memory Architectures
  - Advisor: Prof. Dr. Cevdet Aykanat
- **Bachelor of Science** GPA: 3.44/4.00 September 2007 – June 2012

## RESEARCH EXPERIENCE

---

### Computer Systems Engineer 2 November 2025 – *present* Lawrence Berkeley National Laboratory Berkeley, California, USA

- \* Sparse Symmetric Matrix Factorization

### Postdoctoral Scholar November 2022 – November 2025 Lawrence Berkeley National Laboratory Berkeley, California, USA

- \* Sparse Symmetric Matrix Factorization
- Right-looking sparse  $A = LDL^T$  factorization 2023
  - Developed efficient pivoting strategies for factoring sparse general symmetric matrices
  - Achieved higher accuracy and reduced fill-in
- Right-looking sparse Cholesky factorization 2023
  - Accelerated a serial Cholesky algorithm on GPUs, achieving  $4\times$  speedup; results published in SC24-W
  - Developed a serial right-looking supernodal sparse Cholesky factorization
  - Obtained a Cholesky algorithm that requires no temporary working storage and avoids assembly operations

### Postdoctoral Researcher September 2020 – October 2022 Bilkent University Ankara, Turkey

- \* Stochastic Gradient Descent (SGD) for Matrix Completion
- Scaling stratified SGD for distributed matrix completion 2022
  - Collaborated on a point-to-point (P2P) communication scheme and a related hold-and-combine algorithm
  - Implemented a computational load balancing method and an HP model that minimizes the communication volume
  - Obtained up to  $15\times$  faster parallel runtime and published the results in IEEE TKDE
- Reducing stale data usage and bandwidth requirements in SGD 2022
  - Contributed to the design of the MPI-based SGD with sub-iterations
  - Designed and implemented an HP model that minimizes both staleness and bandwidth requirements in SGD
  - Obtained up to 34% reduction in parallel runtime and published the results in IEEE TC
- Hybrid Parallel SGD 2022
  - Contributed to the design and implementation of the MPI + POSIX-threads-based hybrid parallel SGD
  - Obtained up to  $6\times$  better throughput and published the results in KNOSYS

- \* Parallel Sparse Tensor Decomposition
  - Hiding latency of sparse point-to-point (P2P) communications within dense all-reduce operations 2021
    - Collaborated on reorganizing the CPD algorithm for embedding P2P messages into ALLREDUCE messages
    - Designed and implemented an HP model that minimizes the concurrent communication volume in the embedded ALLREDUCE
    - Scaled CPD on up to 4096 processors and published the results in IEEE TPDS
  - General medium-grain sparse tensor partitioning for distributed a CPD 2019
    - Designed and implemented an HP model for medium-grain partitioning without any topological constraint
    - Utilized the RB paradigm to boost performance at each level of the partitioning
    - Conducted experiments with 10 real-world tensors on 1024 cores and published the results in IEEE TPDS
  - Locality-aware fiber and slice reordering for shared-memory MTTKRP 2017
    - Implemented an HP model for reordering fibers and/or slices of tensors to increase cache locality during MTTKRP
    - Adopted SPLATT's OpenMP-based MTTKRP and conducted experiments
- \* Large Scale Benchmarking
  - Code owner of the NEMO package in the PRACE-6IP T7.4.A activity 2021
    - Prepared architecture specific build files for both NEMO and XIOS packages for six tier-0 HPC systems
    - Benchmarked the NEMO and XIOS on HAWK, TGCC Joliot Curie, JUWELS, Marconi100, MareNostrum, and SuperMUC-NG using up to 10,000 cores
    - Contributed to the deliverable PRACE-6IP-D7.4: Evaluation of Benchmark Performance
- \* Other Hypergraph Partitioning Models
  - Simultaneous computational and data load balancing of the processors on distributed-memory setting 2022
    - Collaborated on the design of two-constraint HP models that encode computational and data load simultaneously
    - Collaborated on experiments with two different applications and published the results in SIAM J. Sci. Comput.

**Teaching and Research Assistant** September 2012 – June 2020  
 Bilkent University Ankara, Turkey

- \* Sparse Matrix Vector Multiplication (SpMV) and Sparse Matrix Dense Matrix Multiplication (SpMM)
  - Volume balancing and latency minimization for reduce operations 2018
    - Formulated a novel vertex weighting scheme for the HP model that balances volume loads across processors
    - Implemented a refinement algorithm invoked during recursive bisection (RB) and decreased the increase in communication volume
    - Performed extensive experiments on 512 processors for 70 matrices
    - Obtained a 30% faster parallel runtime in column-parallel SpMV and published in IEEE TPDS
  - Locality-aware shared-memory parallel  $y \leftarrow AA^T x$  on many-core processors 2014
    - Implemented OpenMP-based  $y \leftarrow AA^T x$  which achieves reuse of A-matrix nonzeros and vector entries
    - Conducted detailed experiments on Intel Xeon Phi co-processor running in offload mode
    - Obtained a 20% reduction in parallel runtime and published in IEEE TPDS
- \* Generalized Sparse Matrix Matrix Multiplication (SpGEMM)
  - Efficient Vectorization of SpGEMM 2016
    - Transformed a  $C = ADB$  instance into  $C = Zd$  SpMV instance by multiplying A-matrix with B-matrix columns
    - Implemented an efficiently vectorized OpenMP-based  $C = Zd$  using AVX instructions
    - Conducted experiments on the Intel Xeon Phi co-processor and Xeon processor

## TECHNICAL SKILLS

---

**Advanced in:** C, C++, MPI, OpenMP, MATLAB, MAGMA

**Familiar with:** CUDA, Fortran, Python, Java SE, C#, Assembly (MIPS and Intel 8051), PHP, SQL

## PAPERS UNDER PREPARATION

---

- **M. O. Karsavuran**, E. G. Ng, B. W. Peyton, "A comparison of two effective methods for reordering columns within supernodes", *arXiv preprint*, 2024. doi: 10.48550/arXiv.2501.08395.
- **M. O. Karsavuran**, E. G. Ng, B. W. Peyton, J. Peyton, "Some new techniques to use in serial sparse Cholesky factorization algorithms", *arXiv preprint*, 2024. doi: 10.48550/arXiv.2409.13090.

## JOURNAL PUBLICATIONS

---

- K. Büyükkaya, **M. O. Karsavuran** and C. Aykanat, "Stochastic Gradient Descent for Matrix Completion: Hybrid Parallelization on Shared- and Distributed-Memory Systems" in *Knowledge-Based Systems*, vol. 283, pp. 111176, Jan. 2024. doi: 10.1016/j.knosys.2023.111176
- N. Abubaker, O. Çağlayan, **M. O. Karsavuran** and C. Aykanat, "Minimizing Staleness and Communication Overhead in Distributed SGD for Collaborative Filtering" in *IEEE Transactions on Computers*, vol. 72, no. 10, pp. 2925-2937, Oct. 2023. doi: 10.1109/TC.2023.3275107
- N. Abubaker, **M. O. Karsavuran** and C. Aykanat, "Scaling Stratified Stochastic Gradient Descent for Distributed Matrix Completion," in *IEEE Transactions on Knowledge and Data Engineering*, vol. 35, no. 10, pp. 10603-10615, Oct. 2023. doi: 10.1109/TKDE.2023.3253791
- M.F. Çeliktug, **M. O. Karsavuran**, S. Acer and C. Aykanat, "Simultaneous Computational and Data Load Balancing in Distributed-Memory Setting," in *SIAM Journal on Scientific Computing*, vol. 44, no. 6, pp. C399-C424, Nov. 2022. doi: 10.1137/22M1485772
- N. Abubaker, **M. O. Karsavuran** and C. Aykanat, "Scalable Unsupervised ML: Latency Hiding in Distributed Sparse Tensor Decomposition," in *IEEE Transactions on Parallel and Distributed Systems*, vol. 33, no. 11, pp. 3028-3040, Nov. 2022. doi: 10.1109/TPDS.2021.3128827
- **M. O. Karsavuran**, S. Acer and C. Aykanat, "Partitioning Models for General Medium-Grain Parallel Sparse Tensor Decomposition," in *IEEE Transactions on Parallel and Distributed Systems*, vol. 32, no. 1, pp. 147-159, Jan. 2021. doi: 10.1109/TPDS.2020.3012624
- **M. O. Karsavuran**, S. Acer and C. Aykanat, "Reduce Operations: Send Volume Balancing While Minimizing Latency," in *IEEE Transactions on Parallel and Distributed Systems*, vol. 31, no. 6, pp. 1461-1473, June 2020. doi: 10.1109/TPDS.2020.2964536
- **M. O. Karsavuran**, K. Akbudak and C. Aykanat, "Locality-Aware Parallel Sparse Matrix-Vector and Matrix-Transpose-Vector Multiplication on Many-Core Processors," in *IEEE Transactions on Parallel and Distributed Systems*, vol. 27, no. 6, pp. 1713-1726, June 2016. doi: 10.1109/TPDS.2015.2453970

## CONFERENCE PUBLICATIONS

---

- **M. O. Karsavuran**, E. G. Ng, B. W. Peyton, "GPU Accelerated Sparse Cholesky Factorization" in *SC24-W: Workshops of the International Conference for High Performance Computing, Networking, Storage and Analysis*, Atlanta, GA, USA, 2024, pp. 703-707. doi: 10.1109/SCW63240.2024.00098

## TALKS & POSTERS

---

- Sparse Cholesky Factorization Utilizing GPUs, SIAM CSE25, TX, USA.
- Heuristics for Robust Factorization of Sparse Symmetric Indefinite Matrices, SIAM LA24, Paris, France.
- Sparse Tensor Partitioning for Scalable Distributed CPD-ALS, SIAM PP22, Virtual.
- Medium-Grain Partitioning for Sparse Tensor Decomposition, SIAM CSE21, Virtual.
- Exploiting Matrix Reuse and Data Locality in Sparse Matrix-Vector and Matrix-Transpose-Vector Multiplication on Many-Core Architectures, SIAM CSC16, NM, USA.

## PROFESSIONAL SERVICE

---

- **Reviewer** for IEEE Transactions on Parallel and Distributed Systems
- **Reviewer** for ACM Transactions on Architecture and Code Optimization
- **Reviewer** for The Journal of Supercomputing
- **Reviewer** for CCPE (Concurrency and Computation: Practice and Experience)
- **Reviewer** for PPAM22 (14th International Conference on Parallel Processing and Applied Mathematics)
- **Reviewer** for BAŞARIM 2020 (6th National Conference on High Performance Computing)
- **Reviewer** for IPDPS 2018 (31st IEEE International Parallel & Distributed Processing Symposium)